

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso, Ridge
```

```
In [2]: df=pd.read_csv(r"C:\Users\DELL E5490\Downloads\Advertising.csv")
df
```

Out[2]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

```
In [3]: df=df[['Sales','Radio','TV','Newspaper']]
df.columns=['sales','radio','tv','newspaper']
```

```
In [4]: df.head()
```

```
Out[4]:
```

	sales	radio	tv	newspaper
0	22.1	37.8	230.1	69.2
1	10.4	39.3	44.5	45.1
2	12.0	45.9	17.2	69.3
3	16.5	41.3	151.5	58.5
4	17.9	10.8	180.8	58.4

```
In [5]: df.describe()
```

```
Out[5]:
```

	sales	radio	tv	newspaper
count	200.000000	200.000000	200.000000	200.000000
mean	15.130500	23.264000	147.042500	30.554000
std	5.283892	14.846809	85.854236	21.778621
min	1.600000	0.000000	0.700000	0.300000
25%	11.000000	9.975000	74.375000	12.750000
50%	16.000000	22.900000	149.750000	25.750000
75%	19.050000	36.525000	218.825000	45.100000
max	27.000000	49.600000	296.400000	114.000000


In [6]: `df.tail()`

Out[6]:

	sales	radio	tv	newspaper
195	7.6	3.7	38.2	13.8
196	14.0	4.9	94.2	8.1
197	14.8	9.3	177.0	6.4
198	25.5	42.0	283.6	66.2
199	18.4	8.6	232.1	8.7

```
In [7]: sns.lmplot(x="sales",y="radio",data=df,order=2,ci=None)
```

```
Out[7]: <seaborn.axisgrid.FacetGrid at 0x264510f57e0>
```



In [8]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   sales       200 non-null    float64
1   radio       200 non-null    float64
2   tv          200 non-null    float64
3   newspaper   200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```


In [9]: df.fillna(method='ffill',inplace=True)
X=np.array(df['sales']).reshape(-1,1)
y=np.array(df['radio']).reshape(-1,1)
df.dropna(inplace=True)

In [10]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)

In [11]: regr=LinearRegression()
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))


0.04288065646388528

```
In [12]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='r')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [13]: df500=df[:][:500]  
  
sns.lmplot(x="sales",y="radio",data=df500,order=1,ci=None)
```


Out[13]: <seaborn.axisgrid.FacetGrid at 0x264519a5300>



In []:

```
In [14]: df500.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
regr=LinearRegression()
regr.fit(X_train,y_train)
print("Regression:",regr.score(X_test,y_test))
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='r')
plt.plot(X_test,y_pred,color='b')
plt.show()
```

Regression: 0.11402599094377575

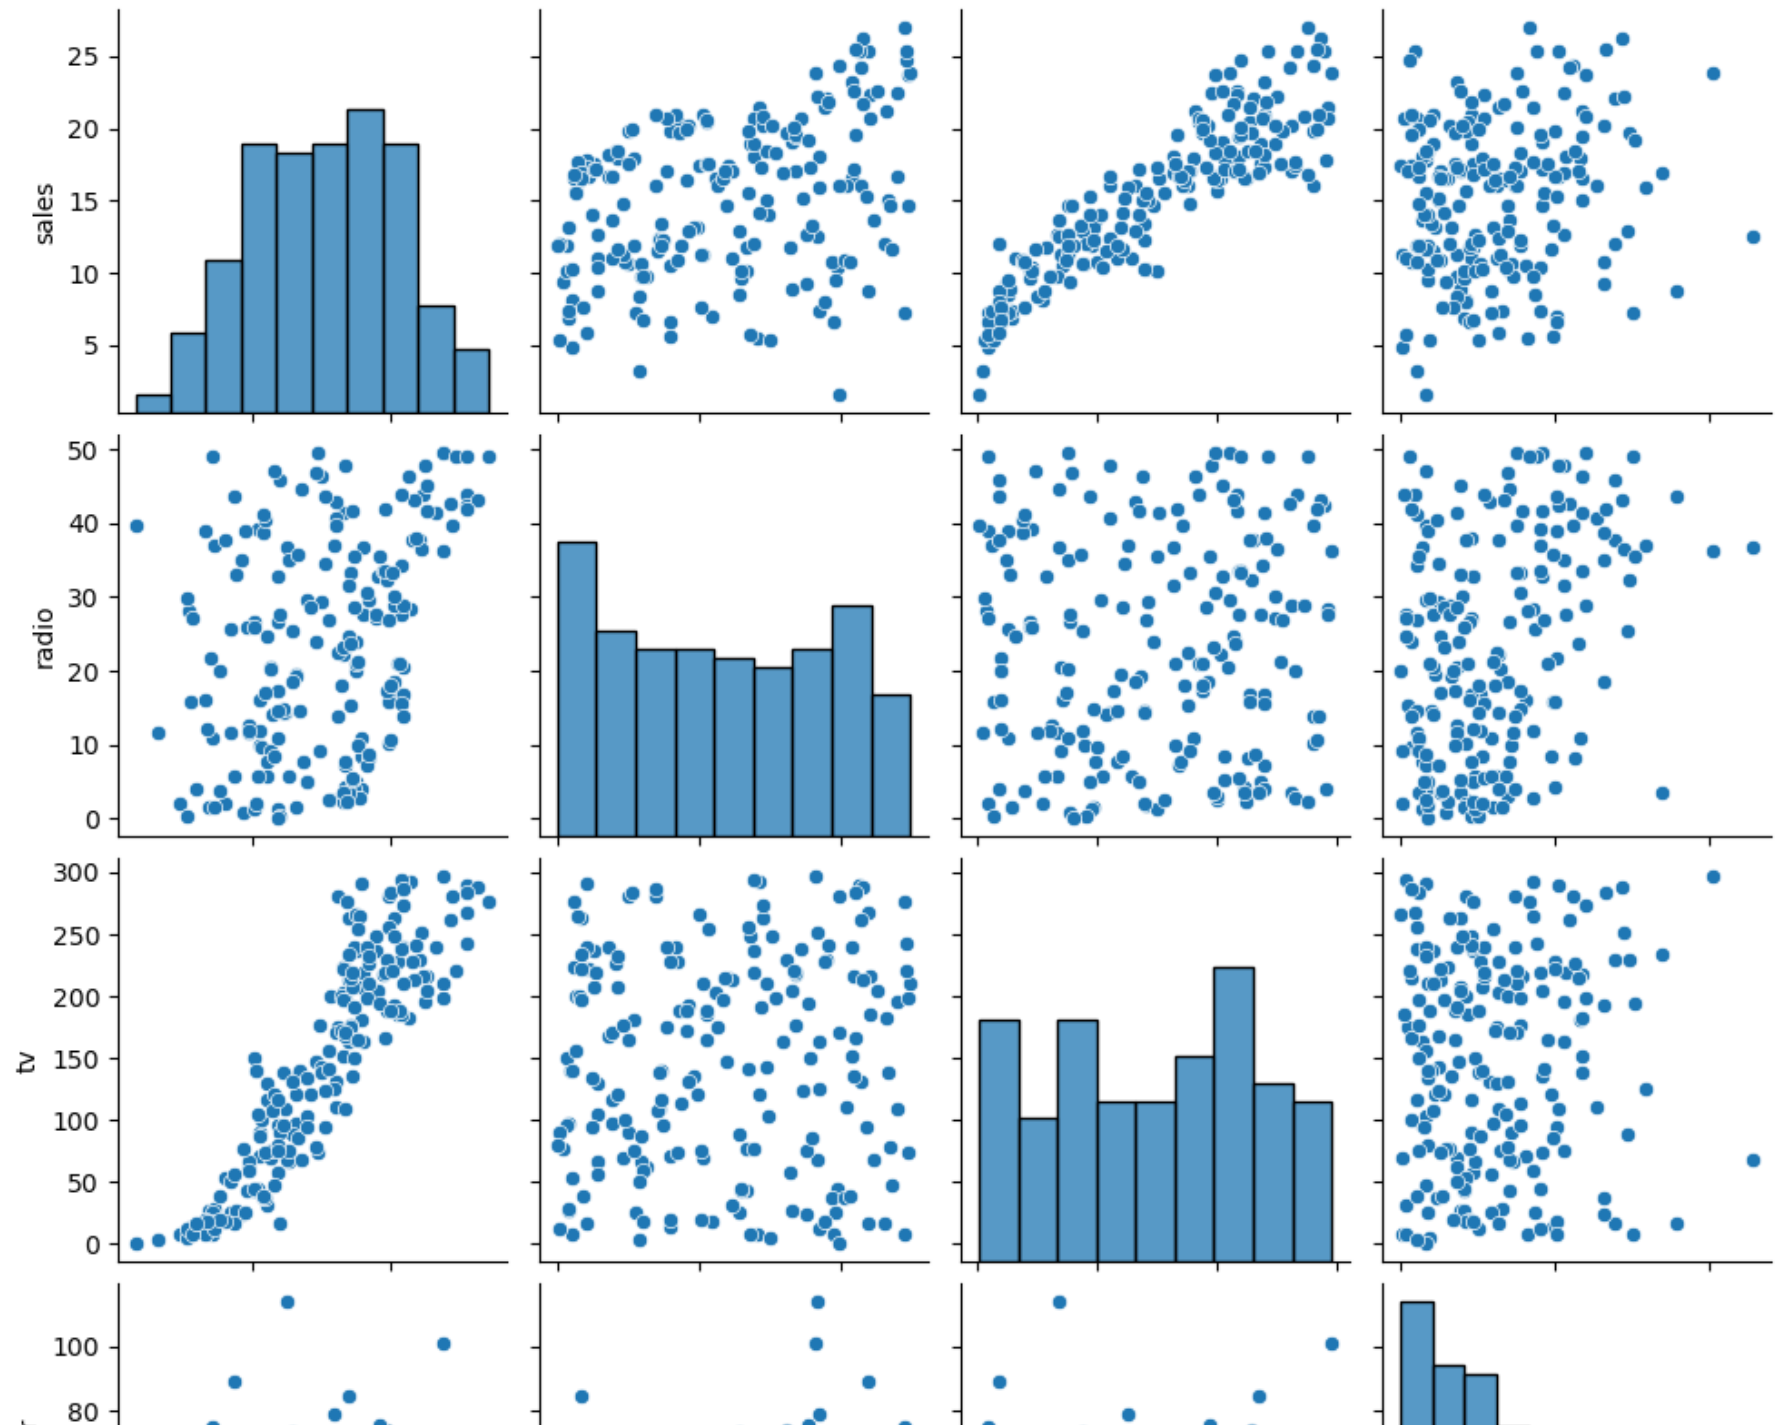




```
In [15]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
#train the model
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2.score:",r2)
```

R2.score: 0.11402599094377575

```
In [16]: sns.pairplot(df)
```


```
Out[16]: <seaborn.axisgrid.PairGrid at 0x264519f7760>
```



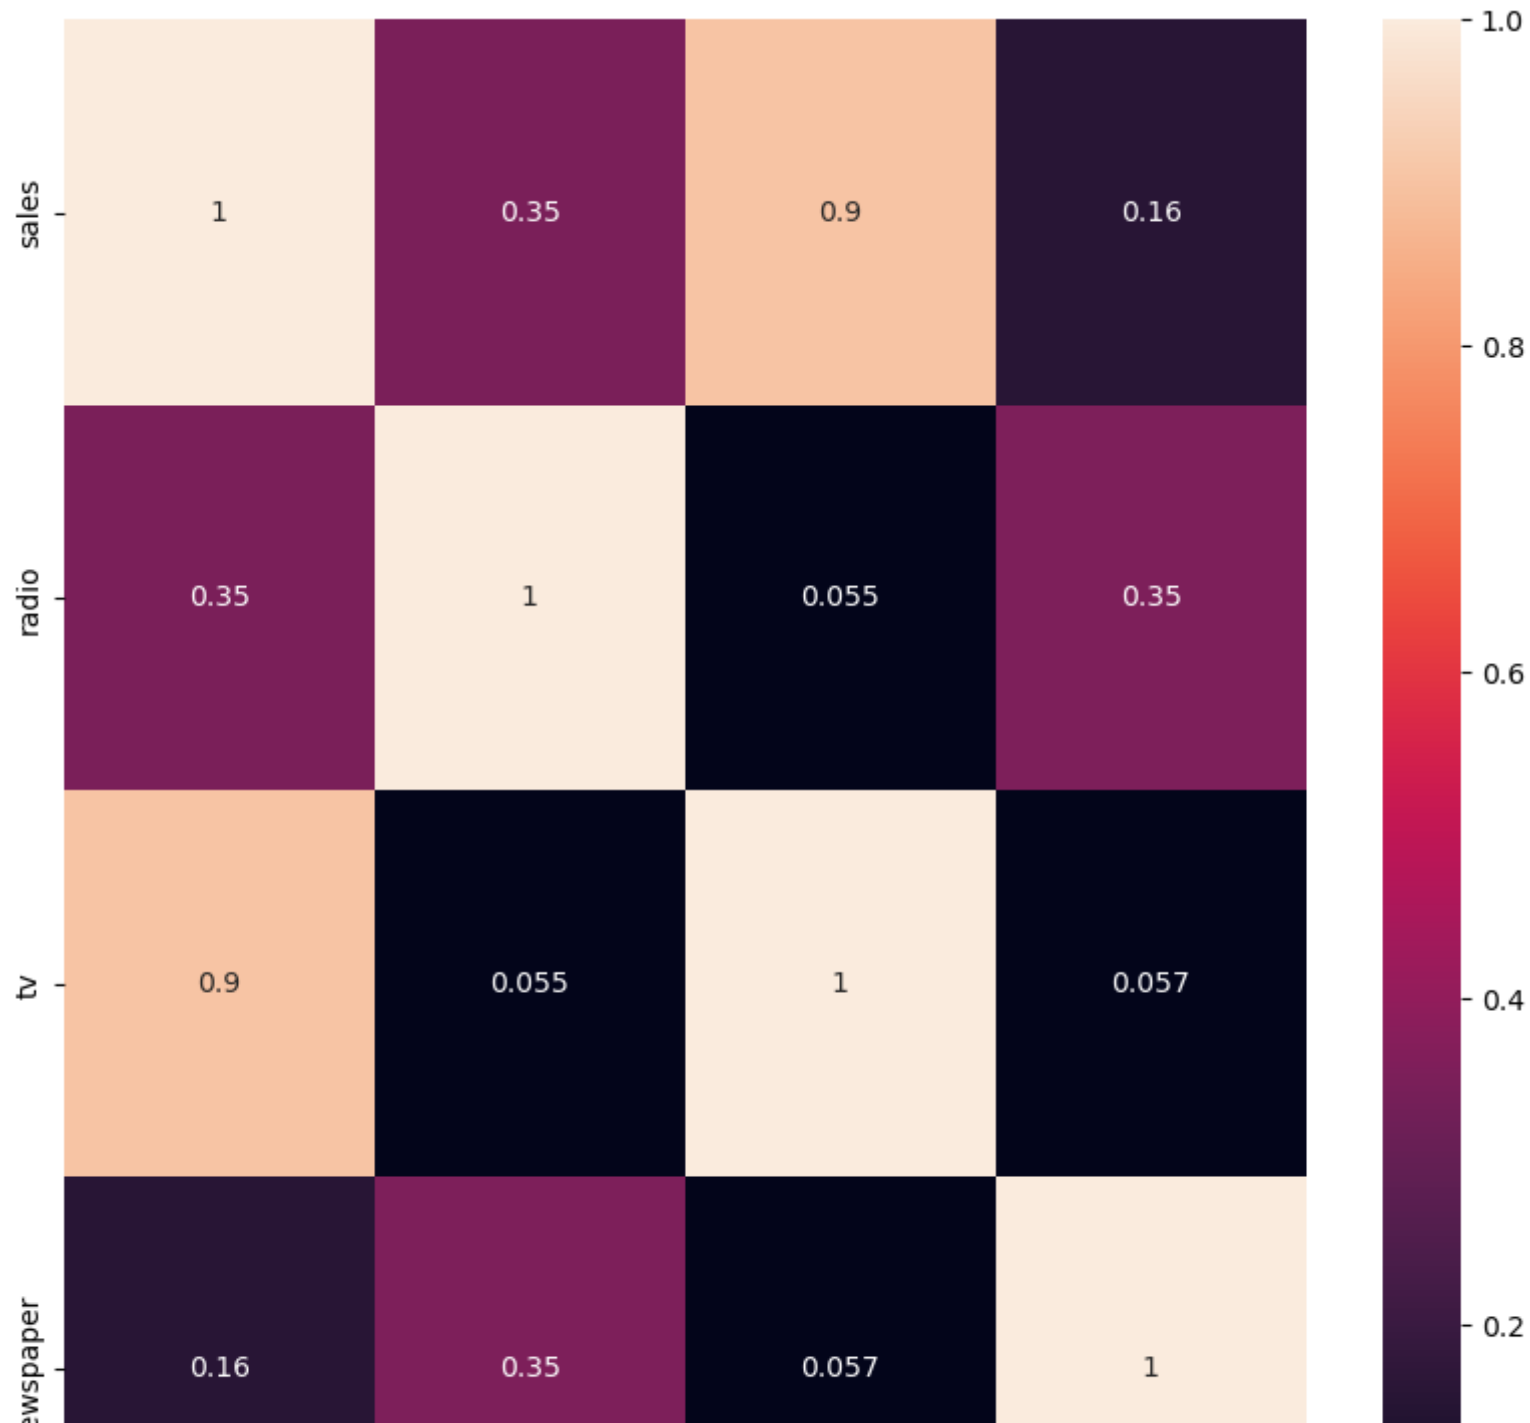
```
In [17]: sns.pairplot(df,x_vars=['tv','radio','newspaper'],y_vars='sales',height=7,aspect=0.7,kind='reg')
```

```
Out[17]: <seaborn.axisgrid.PairGrid at 0x26453636c20>
```



```
In [18]: plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot = True)
```

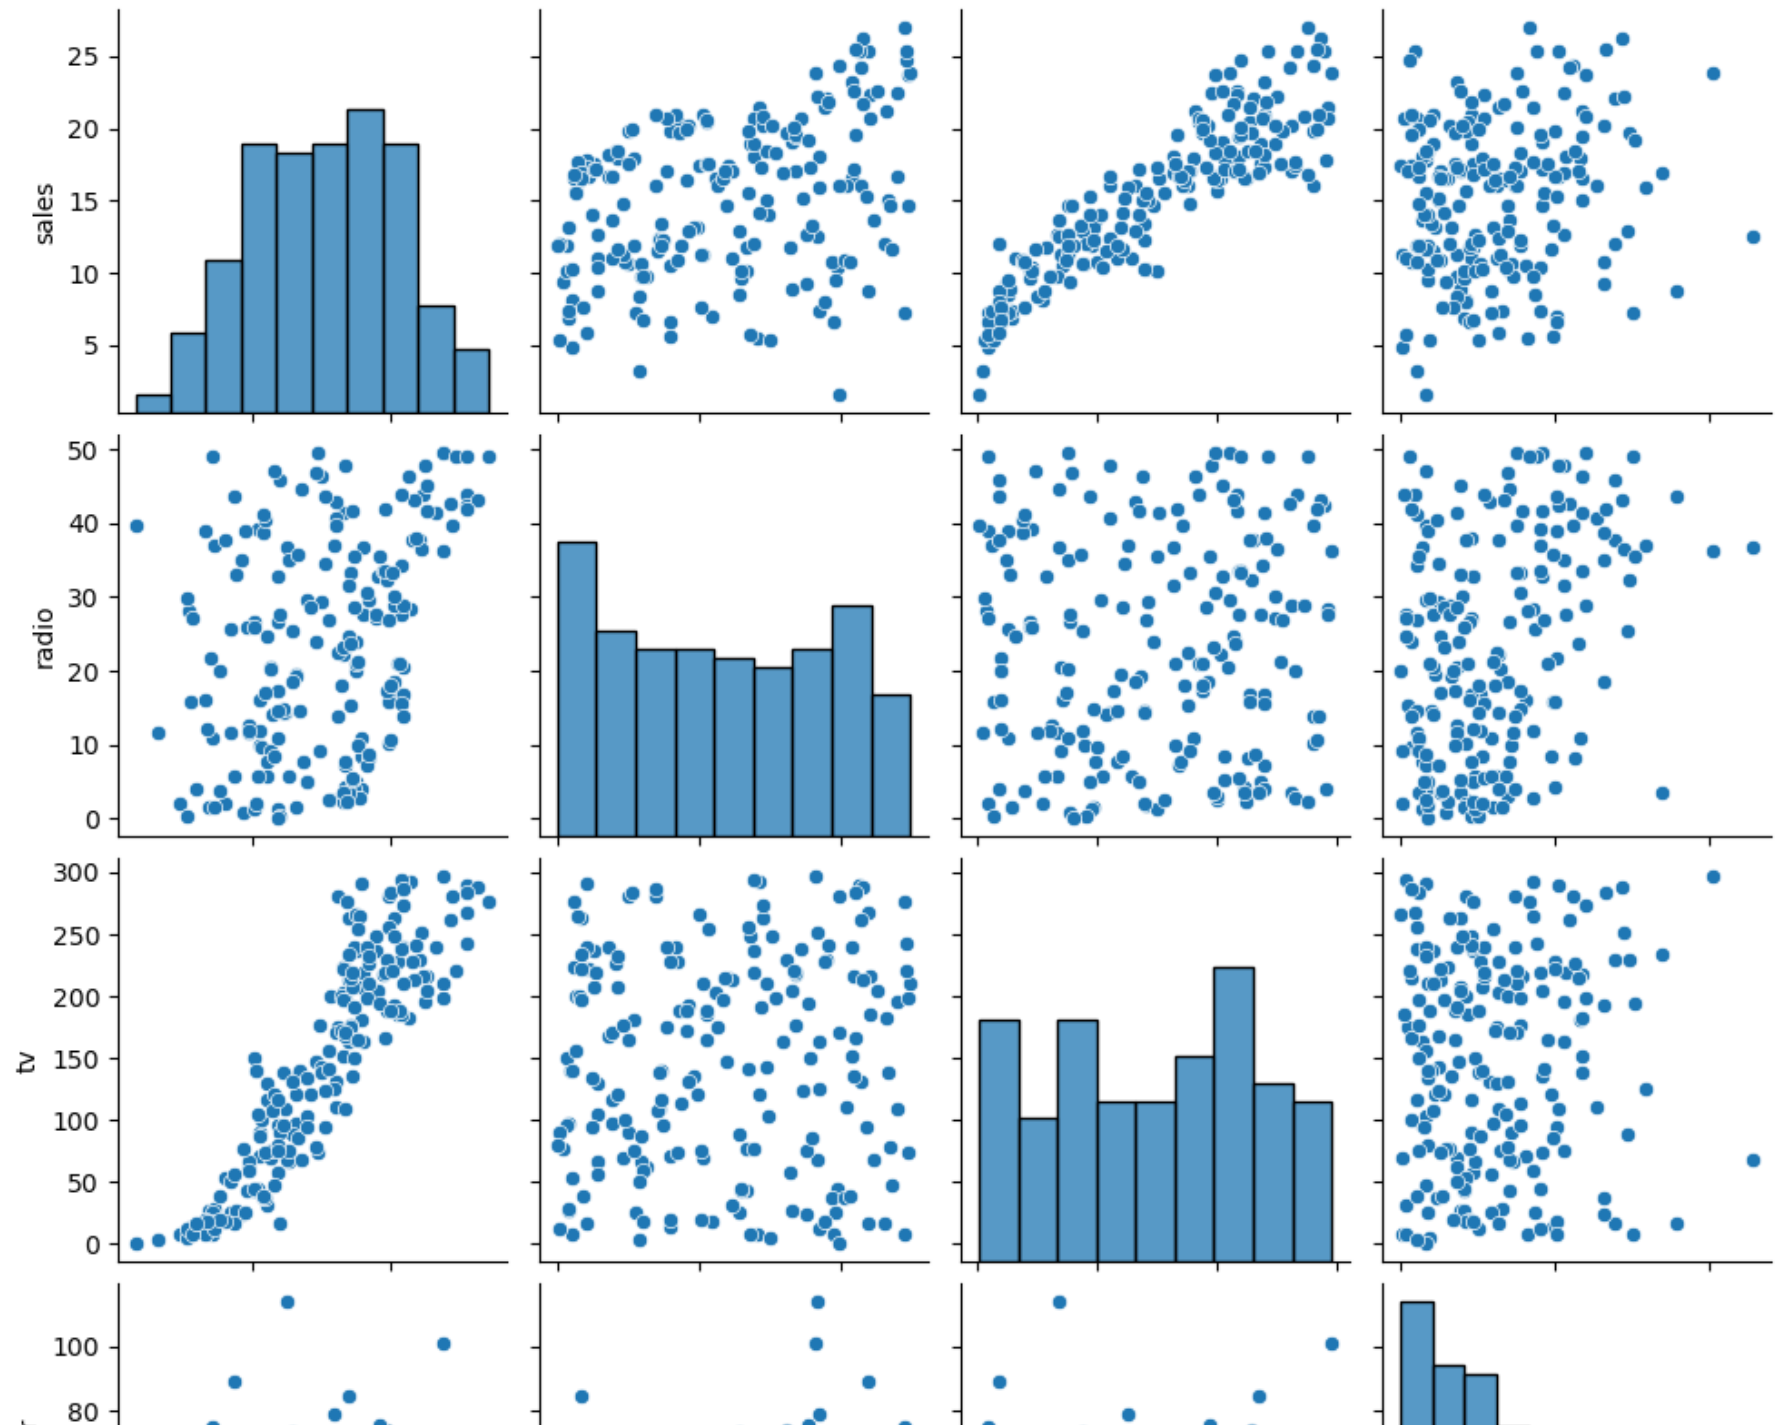
```
Out[18]: <Axes: >
```







In [19]:


```
sns.pairplot(df)  
df.sales = np.log(df.sales)
```







In [20]:

```
#pairplot  
sns.pairplot(df)  
df.sales=np.log(df.sales)
```



```
In [21]: print(regr.score(X_test,y_test))
```

```
0.11402599094377575
```

```
In [22]: features=df.columns[0:2]
target=df.columns[-1]
#X and y values
X=df[features].values
y=df[target].values

#split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3, random_state=17)

print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#scale features
scaler= StandardScaler()
X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
```

```
The dimension of X_train is {} (140, 2)
```

```
The dimension of X_test is {} (60, 2)
```

```
In [23]: #model
lr=LinearRegression()
#fitmodel
lr.fit(X_train,y_train)
#actual
actual=y_test

train_score_lr=lr.score(X_train,y_train)
test_score_lr=lr.score(X_test,y_test)

print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.08495398635901708

The test score for lr model is 0.19849660346463094

```
In [24]: #ridge regression model
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge model:

The train score for ridge model is 0.08462956788927811

The test score for ridge model is 0.18990201599412504

```
In [25]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;$\alpha$')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```






```
In [26]: #lasso regression model
print("\nLasso model: \n")
lasso = Lasso(alpha=10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)

print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```


Lasso model:

The train score for ls model is 0.0

The test score for ls model is -0.0003547334659412815

```
In [27]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind= "bar")
```

Out[27]: <Axes: >



```
In [28]: #using the linear cv model  
from sklearn.linear_model import LassoCV  
  
#lasso cross validation  
lasso_cv = LassoCV(alphas = [0.0001,0.001,0.01,0.1,1,10], random_state=0).fit(X_train,y_train)  
  
#score  
print(lasso_cv.score(X_train,y_train))  
print(lasso_cv.score(X_test,y_test))
```

0.08156518528706347

0.1705627246981769


```
In [29]: #plot size
plt.figure(figsize =(10,10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;  $\alpha$ 
#add plot for Lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;  $\alpha$ =grid$'
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='LinearRegression')

#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge,Lasso and Linear regression model")
plt.show()
```






```
In [30]: #Using the linear cv model
from sklearn.linear_model import RidgeCV

#Ridge Cross validation
ridge_cv= RidgeCV(alphas = [0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)

#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train,y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test,y_test)))
```

The train score for ridge model is 0.08462956788927811
The train score for ridge model is 0.18990201599412537

```
In [31]: from sklearn.linear_model import ElasticNet
```

```
In [37]: regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
```

[0. 0.5159749]
18.55035993077493

```
In [33]: y_pred_elastic=regr.predict(X_train)
```

```
In [36]: mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 578.2547454190322

```
In [ ]:
```